

# Presentation Outline: Tweet-Driven Alpha — Sentiment & Correlation Strategies

---

A 6-slide deck covering strategy mechanics, the ticker filtering pipeline, and live backtest results.

---

# Slide 1 — Overview: What Are We Building?

---

**Title:** Tweet-Driven Alpha: Extracting Market Signals from Social Data

**Key points:**

- Hypothesis: public tweets contain statistically meaningful information about near-term price moves
- Two primary signal extraction methods: **Sentiment Analysis** (NLP-based) and **Correlation Discovery** (statistical)
- Supplementary methods: Bag of Words, Embedding similarity, Multi-timeframe scoring, Ensemble
- All methods feed into a unified backtesting engine trading across 13 tickers (SPY, TLT, GLD, etc.)
- Each method is evaluated on a strict out-of-sample test set — the last 30% of all tweet data by time

**Visual suggestion:** high-level system diagram — tweets → signal methods → backtest engine → portfolio equity curve

---

## Slide 2 — Sentiment Strategy: FinBERT-Powered Signal

**Title:** Sentiment Strategy — How a Tweet Becomes a Trade

Every tweet passes through four sequential filters before a trade is ever placed.

### Stage 1 — Tweet Cleaning (TweetCleaner)

Raw tweets arrive with HTML markup, URLs, and duplicate noise. Before any NLP runs:

| Step                     | What it does  |
|--------------------------|---|
| HTML tag removal         | Strips `<p>`, `<br/>`, etc. via regex                   |
| Entity decoding          | `&` → `&`, `'` → `''                                    |
| URL removal              | Drops all `http(s)://` and `www.` links                 |
| Whitespace normalisation | Collapses multiple spaces, trims                        |
| Short tweet removal      | Drops any tweet with $\leq 5$ characters after cleaning |
| Deduplication            | Removes exact duplicates by `tweet_id` or text          |

### Stage 2 — Importance Filtering (Top 40%)

A composite importance score (0–1) is computed per tweet. **Only the top 40% pass forward.**

| Factor          | Weight | Logic   |
|-----------------|--------|---|
| Length          | 20%    | Longer tweets are more substantive                                    |
| Engagement      | 40%    | replies + reblogs + favourites, normalised                            |
| Keyword density | 30%    | Hits on financial/political terms (tariff, fed, inflation, recession) |
| Uniqueness      | 10%    | Penalises very short or repetitive content                            |

This step eliminates ~60% of raw tweets before any model is run. The effect is fewer but higher-quality signals feeding into FinBERT.

### Stage 3 — FinBERT Sentiment Scoring

The surviving tweets are passed through `ProsusAI/finbert`, a BERT transformer fine-tuned on financial filings and news:

- Outputs three probabilities: **positive, negative, neutral**
- Raw sentiment score = `positive_prob - negative_prob ∈ [-1, +1]`
- Confidence = `abs(positive_prob - negative_prob)` — measures conviction, not just direction
- If the transformer is unavailable, a keyword fallback fires (bullish/bearish word lists)

### Stage 4 — Ticker Relevance Weighting

The same tweet can have different relevance for different tickers. Each ticker has a domain keyword list:

| Ticker | Relevance keywords |
|--------|--------------------|
|--------|--------------------|

|     |  |
|-----|--|
| SPY | market, economy, stocks, equities, s&p |
| TLT | bonds, treasury, interest, fed, rates  |
| GLD | gold, inflation, dollar, currency      |

A tweet mentioning "fed rate decision" gets higher confidence on TLT than on SPY:

```
adjusted_confidence = base_confidence * (1 + ticker_relevance) [capped at 1.0]
influence_score = sentiment_score * adjusted_confidence * 0.005
```

The 0.005 scalar calibrates output to realistic tweet price impact ( $\pm 0.5\%$  per trade).

## Stage 5 — Influence Score Threshold

Only tweets where  $|influence\_score| \geq 0.0003$  trigger a trade. Tweets that pass cleaning and importance filtering but produce a weak sentiment signal are silently discarded at this final gate.

**Visual suggestion:** vertical pipeline diagram — Raw tweets → Clean → Importance filter (drop 60%) → FinBERT → Relevance weight → Threshold gate → Trade

---

# Slide 3 — Correlation Strategy: Statistical Feature Discovery

**Title:** Correlation Strategy — Learning What Features Actually Predict Returns

Correlation goes further than sentiment — instead of classifying text, it learns *which measurable properties of a tweet historically predict price moves*, then scores new tweets against those patterns.

## Stage 1 — Same Cleaning & Importance Filter

Identical to Sentiment: HTML cleaning, deduplication, then top 40% by composite importance score. The strategy only ever sees high-quality tweets.

## Stage 2 — 70 / 30 Time-Ordered Train / Test Split

This is the critical difference from sentiment. Before any statistical analysis runs, the dataset is divided **by time**, not randomly:

```
train set = first 70% of tweets (chronological) ← correlations learned here  
test set = last 30% of tweets (chronological) ← backtesting happens here only
```

Why time-ordered? Because random splitting would leak future information into the training set. A tweet from 2024-Q4 appearing in the training fold would allow correlations to be computed using data the model "shouldn't know yet". The time split enforces strict causality.

## Stage 3 — Safe Feature Extraction

Three categories of features are extracted from each tweet. Price-derived features are **explicitly blocked**:

| Category           | Allowed? | Examples   |
|--------------------|----------|--|
| Embedding features | Yes      | `embedding_pca_0` ... `embedding_pca_49` — semantic co         |
| Engagement metrics | Yes      | `replies_count`, `reblogs_count`, `favourites_count`, `log_to  |
| Tweet-derived      | Yes      | `tweet_word_count`, engagement ratios                          |
| Price features     | **No**   | `ema_*`, `rsi_*`, `macd_*`, `vol_*`, `bb_*`, `atr_*` — blocked |

Price features are blocked because they are computed from bars that are timestamped *after* the tweet, meaning including them would mean trading on information that wasn't available at signal time.

## Stage 4 — Per-Feature Statistical Filtering

On the training set only, for every (feature, ticker) pair:

1. Align tweet feature values to that ticker's 5-minute forward return
2. Compute **Pearson correlation** and **p-value**
3. Compute **mutual information** to catch non-linear relationships
4. Apply filter: keep if  $p\_value < 0.05$  AND  $|correlation| \geq 0.02$
5. If fewer than 3 features survive, relax to top 10 by p-value with  $\min\_abs\_corr \times 0.5$
6. Retain top 20 features per ticker, ranked by  $|correlation|$

The result: a small, statistically significant feature set per ticker. For SPY this might be "high engagement + specific PCA embedding dimensions". For GLD it might be "reblogs count + a different embedding axis".

## Stage 5 — Scoring a New Tweet (Test Set)

For each incoming tweet in the 30% test window:

```
influence_score = Σ (feature_value × correlation_coeff × |weight|)  
÷ sqrt(n_matched_features)
```

The `sqrt(n)` dilution is intentional: dividing by `n` would overly punish tweets with many weak features, while dividing by `sqrt(n)` preserves the signal from a single very strong feature match.

Confidence = `min(1.0, n_matched_features / 10.0)` — a tweet matching all 10 top features gets confidence 1.0.

## Stage 6 — Influence Score Threshold

Same gate as sentiment but tighter:  $|influence\_score| \geq 0.0002$ . The correlation method fires fewer trades (26 vs 55 for sentiment) precisely because its threshold requires a statistically grounded signal, not just a positive/negative classification.

**Visual suggestion:** two-panel diagram — left: training phase (feature selection on 70%), right: inference phase (scoring + threshold on 30%)

---

# Slide 4 — The Full Filtering Pipeline

---

**Title:** From Raw Tweet to Executed Trade — 8 Stages

Both strategies share a common outer pipeline. The differences between them emerge inside stages 4 and 5.

## Pipeline Overview

```
Raw tweet archive (all tickers, all time)
■
[1] Data availability check
■ Only tickers with a .parquet price file survive
■ ~8-10 of 13 tickers pass
↓
[2] Tweet-ticker alignment
■ tweets_df filtered to matched tickers
↓
[3] Tweet cleaning (TweetCleaner)
■ HTML, URLs, dedup, short-tweet removal
↓
[4] Importance filter (top 40%)
■ Composite score: engagement 40%, keywords 30%, length 20%, uniqueness 10%
■ ~60% of tweets dropped here
↓
[5] Method-specific training (70% of data by time)
■ SENTIMENT: FinBERT scores each tweet – no training data needed
■ CORRELATION: Learns per-ticker feature→return correlations on train set
↓
[6] Influence score threshold
■ Sentiment:  $|score| \geq 0.0003$ 
■ Correlation:  $|score| \geq 0.0002$ 
■ Weak signals dropped – no trade fired
↓
[7] Regime-adaptive adjustment (optional)
■ Bull / bear / sideways detection
■ Dynamically raises threshold in low-conviction regimes
↓
[8] Daily trade cap + time-of-day filter
■ Max 3 trades per ticker per day
■ Default: 12:00–14:00 ET only
↓
Trade executed
```

## Where Sentiment and Correlation Diverge

| Stage                | Sentiment                     | Correlation                                    |
|----------------------|-------------------------------|--|
| Training requirement | None — FinBERT is pre-trained | Must learn correlations on 70% train set first |

|                       |  |   |
|-----------------------|--|---|
| Feature input         | Raw tweet text                           | Embedding PCA vectors + engagement numerics     |
| Scoring mechanism     | FinBERT positive/negative probability    | Weighted sum of historical feature correlations |
| Ticker specificity    | Keyword-based relevance boost per ticker | Fully separate feature set per ticker           |
| Influence threshold   | 0.0003                                   | 0.0002  |
| Resulting trade count | 55 (higher volume)                       | 26 (more selective)                             |

## Key Architecture Principle

**Stages 1–5 protect data integrity** — no look-ahead, no circular logic, no price features leaking into tweet scoring.

**Stages 6–8 protect trade quality** — only high-conviction signals, no overtrading, no off-hours noise.

**Visual suggestion:** side-by-side funnel — one column per strategy, tweets narrowing to trades, showing where each strategy diverges

---

## Slide 5 — Results: Equity Curve & Performance Metrics

---

**Title:** Backtest Results — Out-of-Sample Performance (Last 30% of Data)

**Performance table (all methods, out-of-sample):**

| Method          | Total Return | Final Capital | Trades | Win Rate | Avg Return/Trade |
|-----------------|--------------|---------------|--------|----------|------------------|
| **Sentiment**   | **+6.37%**   | \$106,372     | 55     | 58.2%    | 1.02%            |
| **Correlation** | **+4.14%**   | \$104,142     | 26     | 61.5%    | 1.78%            |
| Ensemble        | +3.19%       | \$103,188     | 34     | 55.9%    | 1.37%            |
| Multi-timeframe | +2.57%       | \$102,572     | 19     | 68.4%    | 2.38%            |
| Embedding       | +2.39%       | \$102,388     | 33     | 51.5%    | 1.24%            |
| Bag of Words    | +0.49%       | \$100,492     | 26     | 53.8%    | 0.47%            |

### Key observations:

- Sentiment leads on raw return (+6.37%) with the most trades (55) — high throughput, moderate precision
- Correlation leads on per-trade quality (1.78% avg return, 61.5% win rate) — fewer but more selective trades
- Multi-timeframe has the highest win rate (68.4%) but lowest trade count (19) — very selective filter
- Bag of Words is weakest — proves that surface-level keyword matching without semantic understanding barely works
- All 6 methods beat a flat baseline, validating the hypothesis that tweet signals are exploitable

**Visual suggestion:** overlaid equity curves for all 6 methods from  
[output/tweet\\_ticker\\_methods/all\\_methods\\_comparison.png](output/tweet_ticker_methods/all_methods_comparison.png)

---

## Slide 6 — Risk-Adjusted Returns & Key Takeaways

---

**Title:** Risk-Adjusted Analysis & What We Learned

### Risk-adjusted lens:

- **Sentiment:** highest absolute return but most trades — examine drawdown and Sharpe via [quantstats\\_report.html](#)
- **Correlation:** best return-per-trade ratio (1.78%) — signal is selective and high-conviction by design; the statistical filtering is doing real work
- **Multi-timeframe:** highest win rate (68.4%) with lowest trade count — extremely conservative entry criteria, essentially only trades when multiple time horizons agree
- Core trade-off: *volume* (sentiment, 55 trades) vs *precision* (correlation/multi-timeframe, 19–26 trades)
- Monthly heatmaps available per method:  
[output/tweet\\_ticker\\_methods/{method}/monthly\\_heatmap.png](#)

### Key takeaways:

1. The importance filter (top 40%) is load-bearing — without it, weaker tweets dilute the signal
2. Semantic understanding (FinBERT, embeddings) strongly outperforms surface keyword matching (Bag of Words: +0.49%)
3. The 70/30 time split is not conservative padding — it is the mechanism that prevents the correlation method from cheating
4. The ensemble underperforms the top individual methods — likely signal dilution from weaker sub-strategies being included
5. Next steps: live forward-test correlation and sentiment, explore confidence-scaled position sizing

**Visual suggestion:** bar chart of total return by method + 2x2 quadrant of win rate vs avg return/trade showing precision/volume trade-off

---

## Assets Available for Each Slide

---

All output files are in `output/tweet_ticker_methods/{method}/`:

| File                                      | Use in slide                                 |
|---|--|
| <code>`equity_curve.png`</code>           | Slide 5 — individual equity curves           |
| <code>`equity_curve_simple.png`</code>    | Slide 5 — clean version for presentation     |
| <code>`all_methods_comparison.png`</code> | Slide 5 — overlaid comparison chart          |
| <code>`monthly_heatmap.png`</code>        | Slide 6 — return distribution by month       |
| <code>`metrics.txt`</code>                | Slides 5–6 — source of all numeric metrics   |
| <code>`quantstats_report.html`</code>     | Slide 6 — full Sharpe, Sortino, max drawdown |
| <code>`trades.csv`</code>                 | Slide 4 — trade-level detail if needed       |
| <code>`top_movers_dump.csv`</code>        | Slides 2/3 — examples of top signal tweets   |